

assignment06_FoxAndrea

April 25, 2021

Name: Andrea Fox
Date: April 21, 2021
Assignment: Week 6 Exercises

0.1 Assignment 6.1

Using section 5.1 in Deep Learning with Python as a guide (listing 5.3 in particular), create a

```
[9]: #Load libraries
from keras import layers
from keras import models
from keras.datasets import mnist
from keras.utils import to_categorical
from pathlib import Path
import os
import matplotlib.pyplot as plt
import numpy as np
from keras import losses
from keras import metrics
```

```
[2]: #5.1 Instantiating a small convnet
model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
```

```
[3]: #Looking at the architecture of the convnet
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0

```
conv2d_1 (Conv2D)          (None, 11, 11, 64)          18496
-----
max_pooling2d_1 (MaxPooling2 (None, 5, 5, 64)          0
-----
conv2d_2 (Conv2D)          (None, 3, 3, 64)          36928
=====
Total params: 55,744
Trainable params: 55,744
Non-trainable params: 0
-----
```

```
[4]: #5.2 - Adding a classifier on top of the convnet
model.add(layers.Flatten())
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(10, activation='softmax'))
```

```
[5]: #Checking architecture now to see difference
model.summary()
```

Model: "sequential"

```
-----
Layer (type)                 Output Shape              Param #
=====
conv2d (Conv2D)              (None, 26, 26, 32)       320
-----
max_pooling2d (MaxPooling2D) (None, 13, 13, 32)       0
-----
conv2d_1 (Conv2D)            (None, 11, 11, 64)       18496
-----
max_pooling2d_1 (MaxPooling2 (None, 5, 5, 64)         0
-----
conv2d_2 (Conv2D)            (None, 3, 3, 64)        36928
-----
flatten (Flatten)            (None, 576)              0
-----
dense (Dense)                 (None, 64)               36928
-----
dense_1 (Dense)               (None, 10)               650
=====
Total params: 93,322
Trainable params: 93,322
Non-trainable params: 0
-----
```

```
[21]: #5.3 - Training the convnet on MNIST Images
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()
train_images = train_images.reshape((60000, 28, 28, 1))
```

```

train_images = train_images.astype('float32') / 255

test_images = test_images.reshape((10000, 28, 28, 1))
test_images = test_images.astype('float32') / 255

train_labels = to_categorical(train_labels)
test_labels = to_categorical(test_labels)

model.compile(optimizer='rmsprop',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
history = model.fit(train_images, train_labels, epochs=5, batch_size=64,
                    validation_split=0.1)

```

```

Epoch 1/5
844/844 [=====] - 14s 15ms/step - loss: 0.0151 -
accuracy: 0.9955 - val_loss: 0.0156 - val_accuracy: 0.9960
Epoch 2/5
844/844 [=====] - 12s 15ms/step - loss: 0.0109 -
accuracy: 0.9969 - val_loss: 0.0156 - val_accuracy: 0.9957
Epoch 3/5
844/844 [=====] - 12s 15ms/step - loss: 0.0082 -
accuracy: 0.9976 - val_loss: 0.0218 - val_accuracy: 0.9940
Epoch 4/5
844/844 [=====] - 12s 14ms/step - loss: 0.0079 -
accuracy: 0.9975 - val_loss: 0.0202 - val_accuracy: 0.9958
Epoch 5/5
844/844 [=====] - 12s 14ms/step - loss: 0.0060 -
accuracy: 0.9982 - val_loss: 0.0256 - val_accuracy: 0.9950

```

```

[24]: #Evaluating model on the test data
test_loss, test_acc = model.evaluate(test_images, test_labels)
print('Test Accuracy:', test_acc)
print('Test Loss:', test_loss)

```

```

313/313 [=====] - 1s 5ms/step - loss: 0.0354 -
accuracy: 0.9929
Test Accuracy: 0.992900013923645
Test Loss: 0.035402439534664154

```

```

[29]: #Create plots
history_dict = history.history
acc_values = history_dict['accuracy']
val_acc_values = history_dict['val_accuracy']
val_loss = history_dict['loss']
val_loss_values = history_dict['val_loss']

```


```

epochs = range(1, len(accuracy) + 1)

#Plot for loss
plt.plot(epochs, val_loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss_values, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.show()

```




```

[28]: #Plot for Accuracy
plt.plot(epochs, acc_values, 'bo', label='Training Accuracy')
plt.plot(epochs, val_acc_values, 'b', label = 'Validation Accuracy')
plt.title('Training and Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.show()

```



0.2 Assignment 6.2

Assignment 6.2.a

Using section 5.2 in Deep Learning with Python as a guide, create a ConvNet model that classifies

```
[49]: #load libraries
from keras import layers
from keras import models
from keras import optimizers
from keras.datasets import cifar10
from matplotlib import pyplot
from keras.utils import to_categorical
from pathlib import Path
import os
import matplotlib.pyplot as plt
import numpy as np
from keras import losses
from keras import metrics
```

```
[39]: #5.5 - Instantiating a small convnet for classification
model2 = models.Sequential()
model2.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)))
```

```

model2.add(layers.MaxPooling2D((2, 2)))
model2.add(layers.Conv2D(64, (3, 3), activation='relu'))
model2.add(layers.MaxPooling2D((2, 2)))
model2.add(layers.Conv2D(128, (3, 3), activation='relu'))

```

```

[40]: #Looking at the architecture of the convnet
model2.summary()

```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
conv2d_18 (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d_15 (MaxPooling)	(None, 15, 15, 32)	0
conv2d_19 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_16 (MaxPooling)	(None, 6, 6, 64)	0
conv2d_20 (Conv2D)	(None, 4, 4, 128)	73856

Total params: 93,248
 Trainable params: 93,248
 Non-trainable params: 0

```

[43]: #Add classifier to convnet
model2.add(layers.Flatten())
model2.add(layers.Dense(512, activation='relu'))
model2.add(layers.Dense(10, activation='softmax'))

```

```

[44]: #Second look at architecture to see difference
model2.summary()

```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
conv2d_18 (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d_15 (MaxPooling)	(None, 15, 15, 32)	0
conv2d_19 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_16 (MaxPooling)	(None, 6, 6, 64)	0

conv2d_20 (Conv2D)	(None, 4, 4, 128)	73856

flatten_4 (Flatten)	(None, 2048)	0

dense_8 (Dense)	(None, 512)	1049088

dense_9 (Dense)	(None, 10)	5130
=====		
Total params: 1,147,466		
Trainable params: 1,147,466		
Non-trainable params: 0		

```
[45]: #Training the convnet on cifar10 Images
(train_images2, train_labels2), (test_images2, test_labels2) = cifar10.
↳load_data()

#Found this on general chat from Sam Loyd
train_images2 = train_images2.reshape((50000, 32, 32, 3))
test_images2 = test_images2.reshape((10000, 32, 32, 3))

train_labels2 = to_categorical(train_labels2)
test_labels2 = to_categorical(test_labels2)
```

```
[46]: #5.6 - configuring the model for training. Also worked with Anna Harvey on this
↳portion
model2.compile(optimizer='rmsprop',
               loss='categorical_crossentropy',
               metrics=['accuracy'])
history2 = model2.fit(train_images2, train_labels2, epochs=20, batch_size=64,
↳validation_data=(test_images2, test_labels2))
```

```
Epoch 1/20
782/782 [=====] - 27s 33ms/step - loss: 5.9022 -
accuracy: 0.2415 - val_loss: 1.5796 - val_accuracy: 0.4406
Epoch 2/20
782/782 [=====] - 26s 33ms/step - loss: 1.3356 -
accuracy: 0.5381 - val_loss: 1.4188 - val_accuracy: 0.5440
Epoch 3/20
782/782 [=====] - 25s 32ms/step - loss: 1.1501 -
accuracy: 0.6117 - val_loss: 1.2844 - val_accuracy: 0.5851
Epoch 4/20
782/782 [=====] - 25s 31ms/step - loss: 1.0808 -
accuracy: 0.6366 - val_loss: 1.4950 - val_accuracy: 0.6158
Epoch 5/20
782/782 [=====] - 24s 31ms/step - loss: 1.0598 -
accuracy: 0.6515 - val_loss: 1.2030 - val_accuracy: 0.5845
Epoch 6/20
```

```

782/782 [=====] - 23s 30ms/step - loss: 1.0304 -
accuracy: 0.6588 - val_loss: 1.1846 - val_accuracy: 0.6151
Epoch 7/20
782/782 [=====] - 24s 30ms/step - loss: 1.0089 -
accuracy: 0.6719 - val_loss: 1.3406 - val_accuracy: 0.6001
Epoch 8/20
782/782 [=====] - 24s 30ms/step - loss: 1.0084 -
accuracy: 0.6737 - val_loss: 1.2869 - val_accuracy: 0.5822
Epoch 9/20
782/782 [=====] - 24s 30ms/step - loss: 1.0139 -
accuracy: 0.6690 - val_loss: 1.4236 - val_accuracy: 0.5944
Epoch 10/20
782/782 [=====] - 24s 31ms/step - loss: 1.0224 -
accuracy: 0.6712 - val_loss: 2.6873 - val_accuracy: 0.4206
Epoch 11/20
782/782 [=====] - 23s 30ms/step - loss: 1.0380 -
accuracy: 0.6694 - val_loss: 1.7243 - val_accuracy: 0.5925
Epoch 12/20
782/782 [=====] - 23s 30ms/step - loss: 1.0256 -
accuracy: 0.6720 - val_loss: 1.5756 - val_accuracy: 0.5621
Epoch 13/20
782/782 [=====] - 23s 30ms/step - loss: 1.0372 -
accuracy: 0.6708 - val_loss: 1.4410 - val_accuracy: 0.5844
Epoch 14/20
782/782 [=====] - 23s 29ms/step - loss: 1.0479 -
accuracy: 0.6645 - val_loss: 1.5604 - val_accuracy: 0.5809
Epoch 15/20
782/782 [=====] - 23s 29ms/step - loss: 1.0521 -
accuracy: 0.6640 - val_loss: 2.1028 - val_accuracy: 0.5229
Epoch 16/20
782/782 [=====] - 23s 29ms/step - loss: 1.0554 -
accuracy: 0.6645 - val_loss: 1.9848 - val_accuracy: 0.5034
Epoch 17/20
782/782 [=====] - 22s 29ms/step - loss: 1.0259 -
accuracy: 0.6728 - val_loss: 1.6156 - val_accuracy: 0.5410
Epoch 18/20
782/782 [=====] - 22s 29ms/step - loss: 1.0733 -
accuracy: 0.6608 - val_loss: 2.1186 - val_accuracy: 0.4374
Epoch 19/20
782/782 [=====] - 23s 29ms/step - loss: 1.0764 -
accuracy: 0.6621 - val_loss: 1.3937 - val_accuracy: 0.5986
Epoch 20/20
782/782 [=====] - 22s 28ms/step - loss: 1.0463 -
accuracy: 0.6691 - val_loss: 1.2201 - val_accuracy: 0.6252

```

```

[72]: #Evaluating model on test data
test_loss2, test_acc2 = model2.evaluate(test_images2, test_labels2)

```



```
print('Test Accuracy:', test_acc2)
print('Test Loss:', test_loss2)
```


313/313 [=====] - 2s 7ms/step - loss: 1.2201 -
accuracy: 0.6252
Test Accuracy: 0.6251999735832214
Test Loss: 1.2201370000839233

```
[52]: #Create plots
history_dict2 = history2.history
acc_values2 = history_dict2['accuracy']
val_acc_values2 = history_dict2['val_accuracy']
val_loss2 = history_dict2['loss']
val_loss_values2 = history_dict2['val_loss']

epochs2 = range(1, len(val_loss2) + 1)


#Plot for loss
plt.plot(epochs2, val_loss2, 'bo', label='Training loss')
plt.plot(epochs2, val_loss_values2, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.show()
```



```
[54]: #Plot for Accuracy
plt.plot(epochs2, acc_values2, 'bo', label='Training Accuracy')
plt.plot(epochs2, val_acc_values2, 'b', label = 'Validation Accuracy')
plt.title('Training and Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.show()
```



Assignment 6.2.b

Using section 5.2 in Deep Learning with Python as a guide, create a ConvNet model that classifies

```
[76]: #load libraries
from keras import layers
from keras import models
from keras import optimizers
from keras.datasets import cifar10
from matplotlib import pyplot
from keras.utils import to_categorical
from pathlib import Path
```

```
import os
import matplotlib.pyplot as plt
import numpy as np
from keras import losses
from keras import metrics
from keras.preprocessing.image import ImageDataGenerator
```

```
[70]: #Instantiating convnet for classification
model3 = models.Sequential()
model3.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)))
model3.add(layers.MaxPooling2D((2, 2)))
model3.add(layers.Dropout(0.2))
model3.add(layers.Conv2D(64, (3, 3), activation='relu'))
model3.add(layers.MaxPooling2D((2, 2)))
model3.add(layers.Dropout(0.2))
model3.add(layers.Conv2D(128, (3, 3), activation='relu'))
```

```
[71]: #look at architecture
model3.summary()
```

Model: "sequential_13"

Layer (type)	Output Shape	Param #
conv2d_36 (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d_32 (MaxPooling)	(None, 15, 15, 32)	0
dropout_7 (Dropout)	(None, 15, 15, 32)	0
conv2d_37 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_33 (MaxPooling)	(None, 6, 6, 64)	0
dropout_8 (Dropout)	(None, 6, 6, 64)	0
conv2d_38 (Conv2D)	(None, 4, 4, 128)	73856
Total params: 93,248		
Trainable params: 93,248		
Non-trainable params: 0		

```
[73]: #Add classifier to convnet
model3.add(layers.Flatten())
model3.add(layers.Dense(512, activation='relu'))
```

```
model3.add(layers.Dense(10, activation='softmax'))
```

```
[74]: #See updated architecture
model3.summary()
```

Model: "sequential_13"

Layer (type)	Output Shape	Param #
conv2d_36 (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d_32 (MaxPooling)	(None, 15, 15, 32)	0
dropout_7 (Dropout)	(None, 15, 15, 32)	0
conv2d_37 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_33 (MaxPooling)	(None, 6, 6, 64)	0
dropout_8 (Dropout)	(None, 6, 6, 64)	0
conv2d_38 (Conv2D)	(None, 4, 4, 128)	73856
flatten_5 (Flatten)	(None, 2048)	0
dense_10 (Dense)	(None, 512)	1049088
dense_11 (Dense)	(None, 10)	5130

Total params: 1,147,466
Trainable params: 1,147,466
Non-trainable params: 0

```
[89]: #Training the convnet on cifar10 Images
(train_images3, train_labels3), (test_images3, test_labels3) = cifar10.
    ↳load_data()

#Found this on general chat from Sam Loyd
train_images3 = train_images3.reshape((50000, 32, 32, 3))
test_images3 = test_images3.reshape((10000, 32, 32, 3))

train_labels3 = to_categorical(train_labels3)
test_labels3 = to_categorical(test_labels3)
```

```
[90]: #configuring the model for training. Also worked with Anna Harvey on this
    ↳portion
```

```

model3.compile(optimizer='rmsprop',
               loss='categorical_crossentropy',
               metrics=['accuracy'])
#Got help from Samuel Sears on this portion
train_datagen = ImageDataGenerator(width_shift_range=0.1, height_shift_range=0.
    ↪1, horizontal_flip = True)
train_generator = train_datagen.flow(train_images3, train_labels3, batch_size =
    ↪64)

```

```

[91]: #Samuel Sears said to use 200 epochs as you started seeing better results
    ↪around there and the plots were better
history3 = model3.fit(train_generator, epochs = 200, validation_data =
    ↪(test_images3, test_labels3))

```

```

Epoch 1/200
782/782 [=====] - 54s 68ms/step - loss:
190645433183.1418 - accuracy: 0.1016 - val_loss: 1744569856.0000 - val_accuracy:
0.1000
Epoch 2/200
782/782 [=====] - 53s 67ms/step - loss: 334387327.2031
- accuracy: 0.1025 - val_loss: 2.3098 - val_accuracy: 0.1000
Epoch 3/200
782/782 [=====] - 53s 68ms/step - loss: 2.3290 -
accuracy: 0.1000 - val_loss: 2.3212 - val_accuracy: 0.1000
Epoch 4/200
782/782 [=====] - 53s 68ms/step - loss: 2.3249 -
accuracy: 0.0989 - val_loss: 2.3313 - val_accuracy: 0.1000
Epoch 5/200
782/782 [=====] - 53s 68ms/step - loss: 2.3226 -
accuracy: 0.1024 - val_loss: 2.3266 - val_accuracy: 0.1000
Epoch 6/200
782/782 [=====] - 53s 68ms/step - loss: 2.3230 -
accuracy: 0.0981 - val_loss: 2.3162 - val_accuracy: 0.1000
Epoch 7/200
782/782 [=====] - 53s 67ms/step - loss: 2.3206 -
accuracy: 0.1042 - val_loss: 2.3193 - val_accuracy: 0.1000
Epoch 8/200
782/782 [=====] - 50s 63ms/step - loss: 2.3214 -
accuracy: 0.0967 - val_loss: 2.3196 - val_accuracy: 0.1000
Epoch 9/200
782/782 [=====] - 49s 63ms/step - loss: 2.3209 -
accuracy: 0.0994 - val_loss: 2.3194 - val_accuracy: 0.1000
Epoch 10/200
782/782 [=====] - 49s 63ms/step - loss: 2.3192 -
accuracy: 0.1013 - val_loss: 2.3192 - val_accuracy: 0.1000
Epoch 11/200
782/782 [=====] - 50s 64ms/step - loss: 2.3191 -

```

accuracy: 0.0999 - val_loss: 2.3132 - val_accuracy: 0.1000
Epoch 12/200
782/782 [=====] - 50s 63ms/step - loss: 2.3191 -
accuracy: 0.0988 - val_loss: 2.3344 - val_accuracy: 0.1000
Epoch 13/200
782/782 [=====] - 49s 63ms/step - loss: 2.3189 -
accuracy: 0.0996 - val_loss: 2.3118 - val_accuracy: 0.1000
Epoch 14/200
782/782 [=====] - 49s 63ms/step - loss: 2.3174 -
accuracy: 0.0966 - val_loss: 2.3418 - val_accuracy: 0.1000
Epoch 15/200
782/782 [=====] - 49s 63ms/step - loss: 2.3156 -
accuracy: 0.1055 - val_loss: 2.3137 - val_accuracy: 0.1000
Epoch 16/200
782/782 [=====] - 49s 63ms/step - loss: 2.3171 -
accuracy: 0.0972 - val_loss: 2.3117 - val_accuracy: 0.1000
Epoch 17/200
782/782 [=====] - 49s 63ms/step - loss: 2.3152 -
accuracy: 0.1001 - val_loss: 2.3156 - val_accuracy: 0.1000
Epoch 18/200
782/782 [=====] - 49s 63ms/step - loss: 2.3152 -
accuracy: 0.1013 - val_loss: 2.3172 - val_accuracy: 0.1000
Epoch 19/200
782/782 [=====] - 50s 64ms/step - loss: 2.3149 -
accuracy: 0.0987 - val_loss: 2.3186 - val_accuracy: 0.1000
Epoch 20/200
782/782 [=====] - 50s 63ms/step - loss: 2.3137 -
accuracy: 0.0996 - val_loss: 2.3112 - val_accuracy: 0.1000
Epoch 21/200
782/782 [=====] - 49s 63ms/step - loss: 2.3123 -
accuracy: 0.1018 - val_loss: 2.3140 - val_accuracy: 0.1000
Epoch 22/200
782/782 [=====] - 49s 63ms/step - loss: 2.3134 -
accuracy: 0.0959 - val_loss: 2.3106 - val_accuracy: 0.1000
Epoch 23/200
782/782 [=====] - 50s 63ms/step - loss: 2.3127 -
accuracy: 0.0986 - val_loss: 2.3093 - val_accuracy: 0.1000
Epoch 24/200
782/782 [=====] - 49s 63ms/step - loss: 2.3116 -
accuracy: 0.0984 - val_loss: 2.3052 - val_accuracy: 0.1000
Epoch 25/200
782/782 [=====] - 49s 63ms/step - loss: 2.3110 -
accuracy: 0.1023 - val_loss: 2.3121 - val_accuracy: 0.1000
Epoch 26/200
782/782 [=====] - 49s 63ms/step - loss: 2.3114 -
accuracy: 0.1009 - val_loss: 2.3129 - val_accuracy: 0.1000
Epoch 27/200
782/782 [=====] - 50s 64ms/step - loss: 2.3122 -

accuracy: 0.1007 - val_loss: 2.3058 - val_accuracy: 0.1000
Epoch 28/200
782/782 [=====] - 50s 64ms/step - loss: 2.3102 -
accuracy: 0.0985 - val_loss: 2.3103 - val_accuracy: 0.1000
Epoch 29/200
782/782 [=====] - 50s 64ms/step - loss: 2.3097 -
accuracy: 0.0998 - val_loss: 2.3067 - val_accuracy: 0.1000
Epoch 30/200
782/782 [=====] - 45s 58ms/step - loss: 2.3091 -
accuracy: 0.0993 - val_loss: 2.3076 - val_accuracy: 0.1000
Epoch 31/200
782/782 [=====] - 42s 53ms/step - loss: 2.3078 -
accuracy: 0.1009 - val_loss: 2.3070 - val_accuracy: 0.1000
Epoch 32/200
782/782 [=====] - 48s 61ms/step - loss: 2.3077 -
accuracy: 0.0991 - val_loss: 2.3055 - val_accuracy: 0.1000
Epoch 33/200
782/782 [=====] - 48s 61ms/step - loss: 2.3067 -
accuracy: 0.1003 - val_loss: 2.3060 - val_accuracy: 0.1000
Epoch 34/200
782/782 [=====] - 47s 61ms/step - loss: 2.3065 -
accuracy: 0.1025 - val_loss: 2.3049 - val_accuracy: 0.1000
Epoch 35/200
782/782 [=====] - 47s 61ms/step - loss: 2.3065 -
accuracy: 0.1016 - val_loss: 2.3072 - val_accuracy: 0.1000
Epoch 36/200
782/782 [=====] - 48s 61ms/step - loss: 2.3061 -
accuracy: 0.1007 - val_loss: 2.3056 - val_accuracy: 0.1000
Epoch 37/200
782/782 [=====] - 48s 61ms/step - loss: 2.3056 -
accuracy: 0.0999 - val_loss: 2.3057 - val_accuracy: 0.1000
Epoch 38/200
782/782 [=====] - 48s 61ms/step - loss: 2.3067 -
accuracy: 0.0965 - val_loss: 2.3039 - val_accuracy: 0.1000
Epoch 39/200
782/782 [=====] - 47s 61ms/step - loss: 2.3054 -
accuracy: 0.1009 - val_loss: 2.3040 - val_accuracy: 0.1000
Epoch 40/200
782/782 [=====] - 48s 61ms/step - loss: 2.3047 -
accuracy: 0.0990 - val_loss: 2.3060 - val_accuracy: 0.1000
Epoch 41/200
782/782 [=====] - 48s 61ms/step - loss: 2.3051 -
accuracy: 0.0985 - val_loss: 2.3056 - val_accuracy: 0.1000
Epoch 42/200
782/782 [=====] - 48s 62ms/step - loss: 2.3045 -
accuracy: 0.1021 - val_loss: 2.3033 - val_accuracy: 0.1000
Epoch 43/200
782/782 [=====] - 48s 61ms/step - loss: 2.3044 -

accuracy: 0.1001 - val_loss: 2.3044 - val_accuracy: 0.1000
Epoch 44/200
782/782 [=====] - 47s 60ms/step - loss: 2.3043 -
accuracy: 0.0971 - val_loss: 2.3029 - val_accuracy: 0.1000
Epoch 45/200
782/782 [=====] - 47s 60ms/step - loss: 2.3035 -
accuracy: 0.0978 - val_loss: 2.3038 - val_accuracy: 0.1000
Epoch 46/200
782/782 [=====] - 47s 61ms/step - loss: 2.3045 -
accuracy: 0.0975 - val_loss: 2.3030 - val_accuracy: 0.1000
Epoch 47/200
782/782 [=====] - 47s 60ms/step - loss: 2.3034 -
accuracy: 0.1002 - val_loss: 2.3030 - val_accuracy: 0.1000
Epoch 48/200
782/782 [=====] - 47s 60ms/step - loss: 2.3033 -
accuracy: 0.0988 - val_loss: 2.3035 - val_accuracy: 0.1000
Epoch 49/200
782/782 [=====] - 47s 60ms/step - loss: 2.3035 -
accuracy: 0.0988 - val_loss: 2.3029 - val_accuracy: 0.1000
Epoch 50/200
782/782 [=====] - 47s 61ms/step - loss: 2.3032 -
accuracy: 0.0984 - val_loss: 2.3028 - val_accuracy: 0.1000
Epoch 51/200
782/782 [=====] - 48s 61ms/step - loss: 2.3029 -
accuracy: 0.1010 - val_loss: 2.3029 - val_accuracy: 0.1000
Epoch 52/200
782/782 [=====] - 47s 60ms/step - loss: 2.3031 -
accuracy: 0.0991 - val_loss: 2.3027 - val_accuracy: 0.1000
Epoch 53/200
782/782 [=====] - 51s 65ms/step - loss: 2.3029 -
accuracy: 0.0979 - val_loss: 2.3028 - val_accuracy: 0.1000
Epoch 54/200
782/782 [=====] - 55s 71ms/step - loss: 2.3029 -
accuracy: 0.1000 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 55/200
782/782 [=====] - 55s 71ms/step - loss: 2.3028 -
accuracy: 0.0988 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 56/200
782/782 [=====] - 55s 70ms/step - loss: 2.3028 -
accuracy: 0.0984 - val_loss: 2.3027 - val_accuracy: 0.1000
Epoch 57/200
782/782 [=====] - 55s 70ms/step - loss: 2.3028 -
accuracy: 0.0977 - val_loss: 2.3027 - val_accuracy: 0.1000
Epoch 58/200
782/782 [=====] - 55s 70ms/step - loss: 2.3027 -
accuracy: 0.0998 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 59/200
782/782 [=====] - 55s 71ms/step - loss: 2.3027 -

accuracy: 0.0979 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 60/200
782/782 [=====] - 56s 71ms/step - loss: 2.3027 -
accuracy: 0.0986 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 61/200
782/782 [=====] - 56s 71ms/step - loss: 2.3027 -
accuracy: 0.0986 - val_loss: 2.3027 - val_accuracy: 0.1000
Epoch 62/200
782/782 [=====] - 56s 71ms/step - loss: 2.3027 -
accuracy: 0.0986 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 63/200
782/782 [=====] - 55s 71ms/step - loss: 2.3028 -
accuracy: 0.0992 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 64/200
782/782 [=====] - 55s 70ms/step - loss: 2.3027 -
accuracy: 0.0981 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 65/200
782/782 [=====] - 56s 71ms/step - loss: 2.3027 -
accuracy: 0.1000 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 66/200
782/782 [=====] - 56s 71ms/step - loss: 2.3026 -
accuracy: 0.1020 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 67/200
782/782 [=====] - 51s 66ms/step - loss: 2.3026 -
accuracy: 0.1041 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 68/200
782/782 [=====] - 48s 61ms/step - loss: 2.3027 -
accuracy: 0.0967 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 69/200
782/782 [=====] - 53s 68ms/step - loss: 2.3029 -
accuracy: 0.0975 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 70/200
782/782 [=====] - 55s 70ms/step - loss: 2.3033 -
accuracy: 0.0981 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 71/200
782/782 [=====] - 55s 71ms/step - loss: 2.3032 -
accuracy: 0.0977 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 72/200
782/782 [=====] - 55s 70ms/step - loss: 2.3029 -
accuracy: 0.1000 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 73/200
782/782 [=====] - 55s 70ms/step - loss: 2.3027 -
accuracy: 0.0989 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 74/200
782/782 [=====] - 54s 70ms/step - loss: 2.3027 -
accuracy: 0.1010 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 75/200
782/782 [=====] - 55s 71ms/step - loss: 2.3032 -

accuracy: 0.0968 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 76/200
782/782 [=====] - 55s 70ms/step - loss: 2.3027 -
accuracy: 0.0989 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 77/200
782/782 [=====] - 55s 70ms/step - loss: 2.3029 -
accuracy: 0.0986 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 78/200
782/782 [=====] - 55s 70ms/step - loss: 2.3027 -
accuracy: 0.0994 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 79/200
782/782 [=====] - 55s 70ms/step - loss: 2.3031 -
accuracy: 0.0991 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 80/200
782/782 [=====] - 55s 71ms/step - loss: 2.3027 -
accuracy: 0.1026 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 81/200
782/782 [=====] - 55s 70ms/step - loss: 2.3027 -
accuracy: 0.1002 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 82/200
782/782 [=====] - 55s 70ms/step - loss: 2.3027 -
accuracy: 0.0956 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 83/200
782/782 [=====] - 55s 70ms/step - loss: 2.3029 -
accuracy: 0.0991 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 84/200
782/782 [=====] - 55s 70ms/step - loss: 2.3027 -
accuracy: 0.0966 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 85/200
782/782 [=====] - 52s 67ms/step - loss: 2.3027 -
accuracy: 0.0995 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 86/200
782/782 [=====] - 48s 61ms/step - loss: 2.3027 -
accuracy: 0.1004 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 87/200
782/782 [=====] - 48s 62ms/step - loss: 2.3027 -
accuracy: 0.0992 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 88/200
782/782 [=====] - 48s 61ms/step - loss: 2.3027 -
accuracy: 0.0978 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 89/200
782/782 [=====] - 48s 61ms/step - loss: 2.3026 -
accuracy: 0.0978 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 90/200
782/782 [=====] - 49s 63ms/step - loss: 2.3038 -
accuracy: 0.1004 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 91/200
782/782 [=====] - 48s 61ms/step - loss: 2.3028 -

accuracy: 0.0989 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 92/200
782/782 [=====] - 48s 61ms/step - loss: 2.3027 -
accuracy: 0.0994 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 93/200
782/782 [=====] - 48s 61ms/step - loss: 2.3026 -
accuracy: 0.0978 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 94/200
782/782 [=====] - 48s 61ms/step - loss: 2.3027 -
accuracy: 0.0989 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 95/200
782/782 [=====] - 48s 61ms/step - loss: 2.3028 -
accuracy: 0.0994 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 96/200
782/782 [=====] - 48s 61ms/step - loss: 2.3038 -
accuracy: 0.0998 - val_loss: 2.3027 - val_accuracy: 0.1000
Epoch 97/200
782/782 [=====] - 52s 67ms/step - loss: 2.3028 -
accuracy: 0.0989 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 98/200
782/782 [=====] - 56s 72ms/step - loss: 2.3028 -
accuracy: 0.0970 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 99/200
782/782 [=====] - 57s 73ms/step - loss: 2.3031 -
accuracy: 0.0989 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 100/200
782/782 [=====] - 56s 72ms/step - loss: 2.3027 -
accuracy: 0.0958 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 101/200
782/782 [=====] - 57s 73ms/step - loss: 2.3027 -
accuracy: 0.1009 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 102/200
782/782 [=====] - 58s 74ms/step - loss: 2.3027 -
accuracy: 0.0994 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 103/200
782/782 [=====] - 57s 73ms/step - loss: 2.3028 -
accuracy: 0.0952 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 104/200
782/782 [=====] - 58s 74ms/step - loss: 2.3026 -
accuracy: 0.0986 - val_loss: 2.3027 - val_accuracy: 0.1000
Epoch 105/200
782/782 [=====] - 58s 74ms/step - loss: 2.3029 -
accuracy: 0.0963 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 106/200
782/782 [=====] - 64s 82ms/step - loss: 2.3033 -
accuracy: 0.0994 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 107/200
782/782 [=====] - 57s 73ms/step - loss: 2.3030 -

```

accuracy: 0.0980 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 108/200
782/782 [=====] - 57s 73ms/step - loss: 2.3029 -
accuracy: 0.0974 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 109/200
782/782 [=====] - 58s 74ms/step - loss: 2.3027 -
accuracy: 0.0988 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 110/200
782/782 [=====] - 58s 74ms/step - loss: 2.3027 -
accuracy: 0.1001 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 111/200
782/782 [=====] - 58s 73ms/step - loss: 2.3027 -
accuracy: 0.0958 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 112/200
782/782 [=====] - 57s 73ms/step - loss: 2.3028 -
accuracy: 0.0962 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 113/200
782/782 [=====] - 57s 73ms/step - loss: 2.3088 -
accuracy: 0.0982 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 114/200
782/782 [=====] - 48s 62ms/step - loss: 2.3027 -
accuracy: 0.0968 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 115/200
782/782 [=====] - 48s 61ms/step - loss: 2.3029 -
accuracy: 0.0985 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 116/200
782/782 [=====] - 55s 71ms/step - loss: 2.3027 -
accuracy: 0.0958 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 117/200
782/782 [=====] - 57s 73ms/step - loss: 2.3069 -
accuracy: 0.1006 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 118/200
782/782 [=====] - 58s 74ms/step - loss: 2.3027 -
accuracy: 0.0987 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 119/200
782/782 [=====] - 56s 72ms/step - loss: 2.3028 -
accuracy: 0.0983 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 120/200
782/782 [=====] - 57s 73ms/step - loss: 2.3028 -
accuracy: 0.0996 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 121/200
782/782 [=====] - 57s 73ms/step - loss: 2.3028 -
accuracy: 0.1005 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 122/200
782/782 [=====] - 57s 72ms/step - loss: 2.3027 -
accuracy: 0.0991 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 123/200
782/782 [=====] - 49s 62ms/step - loss: 2.3037 -

```

accuracy: 0.1019 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 124/200
 782/782 [=====] - 48s 62ms/step - loss: 2.3029 -
 accuracy: 0.1012 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 125/200
 782/782 [=====] - 49s 63ms/step - loss: 2.3027 -
 accuracy: 0.1007 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 126/200
 782/782 [=====] - 48s 62ms/step - loss: 2.3027 -
 accuracy: 0.1017 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 127/200
 782/782 [=====] - 48s 62ms/step - loss: 2.3027 -
 accuracy: 0.0960 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 128/200
 782/782 [=====] - 56s 71ms/step - loss: 2.3028 -
 accuracy: 0.1005 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 129/200
 782/782 [=====] - 57s 72ms/step - loss: 2.3027 -
 accuracy: 0.0972 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 130/200
 782/782 [=====] - 57s 73ms/step - loss: 2.3027 -
 accuracy: 0.0974 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 131/200
 782/782 [=====] - 57s 73ms/step - loss: 2.3030 -
 accuracy: 0.0974 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 132/200
 782/782 [=====] - 58s 74ms/step - loss: 2.3027 -
 accuracy: 0.0990 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 133/200
 782/782 [=====] - 59s 75ms/step - loss: 2.3027 -
 accuracy: 0.1003 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 134/200
 782/782 [=====] - 58s 74ms/step - loss: 2.3026 -
 accuracy: 0.0993 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 135/200
 782/782 [=====] - 51s 66ms/step - loss: 2.3027 -
 accuracy: 0.1007 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 136/200
 782/782 [=====] - 56s 72ms/step - loss: 2.3031 -
 accuracy: 0.0972 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 137/200
 782/782 [=====] - 55s 70ms/step - loss: 2.3027 -
 accuracy: 0.0978 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 138/200
 782/782 [=====] - 54s 69ms/step - loss: 2.3027 -
 accuracy: 0.0977 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 139/200
 782/782 [=====] - 53s 68ms/step - loss: 2.3027 -

accuracy: 0.1011 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 140/200
 782/782 [=====] - 54s 69ms/step - loss: 2.3027 -
 accuracy: 0.1006 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 141/200
 782/782 [=====] - 56s 72ms/step - loss: 2.3027 -
 accuracy: 0.0994 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 142/200
 782/782 [=====] - 54s 69ms/step - loss: 2.3027 -
 accuracy: 0.0971 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 143/200
 782/782 [=====] - 54s 69ms/step - loss: 2.3026 -
 accuracy: 0.1009 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 144/200
 782/782 [=====] - 64s 82ms/step - loss: 2.3036 -
 accuracy: 0.1002 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 145/200
 782/782 [=====] - 53s 68ms/step - loss: 2.3028 -
 accuracy: 0.0974 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 146/200
 782/782 [=====] - 53s 68ms/step - loss: 2.3027 -
 accuracy: 0.0994 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 147/200
 782/782 [=====] - 53s 68ms/step - loss: 2.3027 -
 accuracy: 0.1008 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 148/200
 782/782 [=====] - 52s 66ms/step - loss: 2.3027 -
 accuracy: 0.0989 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 149/200
 782/782 [=====] - 50s 63ms/step - loss: 2.3028 -
 accuracy: 0.0949 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 150/200
 782/782 [=====] - 49s 63ms/step - loss: 2.3027 -
 accuracy: 0.0996 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 151/200
 782/782 [=====] - 49s 63ms/step - loss: 2.3027 -
 accuracy: 0.0968 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 152/200
 782/782 [=====] - 50s 64ms/step - loss: 2.3027 -
 accuracy: 0.0977 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 153/200
 782/782 [=====] - 49s 63ms/step - loss: 2.3026 -
 accuracy: 0.1018 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 154/200
 782/782 [=====] - 48s 61ms/step - loss: 2.3027 -
 accuracy: 0.0961 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 155/200
 782/782 [=====] - 46s 59ms/step - loss: 2.3027 -

accuracy: 0.0969 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 156/200
 782/782 [=====] - 46s 59ms/step - loss: 2.3027 -
 accuracy: 0.1006 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 157/200
 782/782 [=====] - 52s 67ms/step - loss: 2.3027 -
 accuracy: 0.1002 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 158/200
 782/782 [=====] - 54s 69ms/step - loss: 2.3027 -
 accuracy: 0.0996 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 159/200
 782/782 [=====] - 55s 70ms/step - loss: 2.3027 -
 accuracy: 0.0992 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 160/200
 782/782 [=====] - 55s 70ms/step - loss: 2.3027 -
 accuracy: 0.1017 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 161/200
 782/782 [=====] - 55s 70ms/step - loss: 2.3027 -
 accuracy: 0.0978 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 162/200
 782/782 [=====] - 60s 77ms/step - loss: 2.3027 -
 accuracy: 0.1002 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 163/200
 782/782 [=====] - 55s 70ms/step - loss: 2.3027 -
 accuracy: 0.0964 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 164/200
 782/782 [=====] - 58s 74ms/step - loss: 2.3050 -
 accuracy: 0.0999 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 165/200
 782/782 [=====] - 64s 82ms/step - loss: 2.3039 -
 accuracy: 0.0978 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 166/200
 782/782 [=====] - 64s 82ms/step - loss: 2.3027 -
 accuracy: 0.0993 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 167/200
 782/782 [=====] - 64s 82ms/step - loss: 2.3027 -
 accuracy: 0.0968 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 168/200
 782/782 [=====] - 64s 82ms/step - loss: 2.3026 -
 accuracy: 0.0963 - val_loss: 2.3027 - val_accuracy: 0.1000
 Epoch 169/200
 782/782 [=====] - 63s 81ms/step - loss: 2.3027 -
 accuracy: 0.1021 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 170/200
 782/782 [=====] - 63s 81ms/step - loss: 2.3034 -
 accuracy: 0.0984 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 171/200
 782/782 [=====] - 63s 80ms/step - loss: 2.3027 -

accuracy: 0.0972 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 172/200
 782/782 [=====] - 63s 80ms/step - loss: 2.3026 -
 accuracy: 0.0988 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 173/200
 782/782 [=====] - 63s 80ms/step - loss: 2.3027 -
 accuracy: 0.1002 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 174/200
 782/782 [=====] - 62s 79ms/step - loss: 2.3027 -
 accuracy: 0.1002 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 175/200
 782/782 [=====] - 62s 80ms/step - loss: 2.3027 -
 accuracy: 0.0993 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 176/200
 782/782 [=====] - 62s 80ms/step - loss: 2.3027 -
 accuracy: 0.0939 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 177/200
 782/782 [=====] - 63s 80ms/step - loss: 2.3025 -
 accuracy: 0.0993 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 178/200
 782/782 [=====] - 62s 79ms/step - loss: 2.3027 -
 accuracy: 0.1018 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 179/200
 782/782 [=====] - 62s 80ms/step - loss: 2.3027 -
 accuracy: 0.1001 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 180/200
 782/782 [=====] - 63s 81ms/step - loss: 2.3027 -
 accuracy: 0.0975 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 181/200
 782/782 [=====] - 63s 80ms/step - loss: 2.3028 -
 accuracy: 0.0952 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 182/200
 782/782 [=====] - 63s 81ms/step - loss: 2.3027 -
 accuracy: 0.0996 - val_loss: 2.3027 - val_accuracy: 0.1000
 Epoch 183/200
 782/782 [=====] - 63s 80ms/step - loss: 2.3028 -
 accuracy: 0.0984 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 184/200
 782/782 [=====] - 63s 80ms/step - loss: 2.3027 -
 accuracy: 0.0989 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 185/200
 782/782 [=====] - 63s 81ms/step - loss: 2.3027 -
 accuracy: 0.0979 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 186/200
 782/782 [=====] - 64s 81ms/step - loss: 2.3026 -
 accuracy: 0.0986 - val_loss: 2.3026 - val_accuracy: 0.1000
 Epoch 187/200
 782/782 [=====] - 63s 81ms/step - loss: 2.3027 -


```

accuracy: 0.0996 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 188/200
782/782 [=====] - 63s 80ms/step - loss: 2.3027 -
accuracy: 0.1007 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 189/200
782/782 [=====] - 64s 81ms/step - loss: 2.3026 -
accuracy: 0.1006 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 190/200
782/782 [=====] - 63s 80ms/step - loss: 2.3027 -
accuracy: 0.1009 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 191/200
782/782 [=====] - 63s 81ms/step - loss: 2.3027 -
accuracy: 0.1002 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 192/200
782/782 [=====] - 63s 80ms/step - loss: 2.3026 -
accuracy: 0.0996 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 193/200
782/782 [=====] - 62s 80ms/step - loss: 2.3027 -
accuracy: 0.0988 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 194/200
782/782 [=====] - 62s 79ms/step - loss: 2.3027 -
accuracy: 0.0992 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 195/200
782/782 [=====] - 58s 74ms/step - loss: 2.3027 -
accuracy: 0.0996 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 196/200
782/782 [=====] - 59s 75ms/step - loss: 2.3026 -
accuracy: 0.0989 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 197/200
782/782 [=====] - 59s 75ms/step - loss: 2.3030 -
accuracy: 0.0995 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 198/200
782/782 [=====] - 59s 75ms/step - loss: 2.3027 -
accuracy: 0.1002 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 199/200
782/782 [=====] - 58s 74ms/step - loss: 2.3024 -
accuracy: 0.1002 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 200/200
782/782 [=====] - 52s 67ms/step - loss: 2.3027 -
accuracy: 0.1006 - val_loss: 2.3026 - val_accuracy: 0.1000

```

```

[92]: #Evaluating model on test data
test_loss3, test_acc3 = model3.evaluate(test_images3, test_labels3)
print('Test Accuracy:', test_acc3)
print('Test Loss:', test_loss3)

```

```

313/313 [=====] - 2s 7ms/step - loss: 2.3026 -
accuracy: 0.1000

```

Test Accuracy: 0.10000000149011612

Test Loss: 2.302603006362915

```
[95]: #Create plots
history_dict3 = history3.history
acc_values3 = history_dict3['accuracy']
val_acc_values3 = history_dict3['val_accuracy']
val_loss3 = history_dict2['loss']
val_loss_values3 = history_dict3['val_loss']

epochs3 = range(1, len(val_loss3) + 1)

#Plot for loss
plt.plot(epochs3, val_loss3, 'bo', label='Training loss')
plt.plot(epochs3, val_loss_values3, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.show()
```

```
-----
ValueError                                Traceback (most recent call last)
<ipython-input-95-e6bce6c06167> in <module>
    10 #Plot for loss
    11 plt.plot(epochs3, val_loss3, 'bo', label='Training loss')
----> 12 plt.plot(epochs3, val_loss_values3, 'b', label='Validation loss')
    13 plt.title('Training and validation loss')
    14 plt.xlabel('Epochs')

/opt/conda/lib/python3.8/site-packages/matplotlib/pyplot.py in plot(scalex, scaley, data, *args, **kwargs)
    2838 @_copy_docstring_and_deprecators(Axes.plot)
    2839 def plot(*args, scalex=True, scaley=True, data=None, **kwargs):
-> 2840     return gca().plot(

    2841         *args, scalex=scalex, scaley=scaley,
    2842         **({"data": data} if data is not None else {}), **kwargs)

/opt/conda/lib/python3.8/site-packages/matplotlib/axes/_axes.py in plot(self, scalex, scaley, data, *args, **kwargs)
    1741         """
    1742         kwargs = cbook.normalize_kwargs(kwargs, mlines.Line2D)
-> 1743         lines = [*self._get_lines(*args, data=data, **kwargs)]
    1744         for line in lines:
    1745             self.add_line(line)
```


```

/opt/conda/lib/python3.8/site-packages/matplotlib/axes/_base.py in
↳ __call__(self, data, *args, **kwargs)
    271         this += args[0],
    272         args = args[1:]
--> 273         yield from self._plot_args(this, kwargs)
    274
    275     def get_next_color(self):

/opt/conda/lib/python3.8/site-packages/matplotlib/axes/_base.py in
↳ _plot_args(self, tup, kwargs)
    397
    398     if x.shape[0] != y.shape[0]:
--> 399         raise ValueError(f"x and y must have same first dimension,
↳ but "
    400                             f"have shapes {x.shape} and {y.shape}")
    401     if x.ndim > 2 or y.ndim > 2:

ValueError: x and y must have same first dimension, but have shapes (20,) and
↳ (200,)

```



```

[96]: #Plot for Accuracy
plt.plot(epochs3, acc_values3, 'bo', label='Training Accuracy')
plt.plot(epochs3, val_acc_values3, 'b', label = 'Validation Accuracy')
plt.title('Training and Validation Accuracy')
plt.xlabel('Epochs')

```

```
plt.ylabel('Accuracy')
plt.legend()

plt.show()
```

```
-----
ValueError                                Traceback (most recent call last)
<ipython-input-96-4890c52fdf6e> in <module>
      1 #Plot for Accuracy
----> 2 plt.plot(epochs3, acc_values3, 'bo', label='Training Accuracy')
      3 plt.plot(epochs3, val_acc_values3, 'b', label = 'Validation Accuracy')
      4 plt.title('Training and Validation Accuracy')
      5 plt.xlabel('Epochs')

/opt/conda/lib/python3.8/site-packages/matplotlib/pyplot.py in plot(scalex,
↳ scaley, data, *args, **kwargs)
    2838 @_copy_docstring_and_deprecators(Axes.plot)
    2839 def plot(*args, scalex=True, scaley=True, data=None, **kwargs):
-> 2840     return gca().plot(

    2841         *args, scalex=scalex, scaley=scaley,
    2842         **({"data": data} if data is not None else {}), **kwargs)


/opt/conda/lib/python3.8/site-packages/matplotlib/axes/_axes.py in plot(self,
↳ scalex, scaley, data, *args, **kwargs)
    1741     """
    1742     kwargs = cbook.normalize_kwargs(kwargs, mlines.Line2D)
-> 1743     lines = [self._get_lines(*args, data=data, **kwargs)]
    1744     for line in lines:
    1745         self.add_line(line)

/opt/conda/lib/python3.8/site-packages/matplotlib/axes/_base.py in
↳ __call__(self, data, *args, **kwargs)
    271         this += args[0],
    272         args = args[1:]
--> 273         yield from self._plot_args(this, kwargs)
    274
    275     def get_next_color(self):

/opt/conda/lib/python3.8/site-packages/matplotlib/axes/_base.py in
↳ _plot_args(self, tup, kwargs)
    397
    398     if x.shape[0] != y.shape[0]:
--> 399         raise ValueError(f"x and y must have same first dimension,
↳ but "

    400                             f"have shapes {x.shape} and {y.shape}")
    401     if x.ndim > 2 or y.ndim > 2:
```

`ValueError: x and y must have same first dimension, but have shapes (20,) and (200,)`



0.3 Assignment 6.3

Load the ResNet50 model. Perform image classification on five to ten images of your choice. The

```
[31]: #load libraries
import os
from pathlib import Path

#load ResNet50 libraries
import numpy as np

from keras.preprocessing.image import img_to_array
from keras.applications.resnet50 import preprocess_input
from keras.applications.imagenet_utils import decode_predictions
import matplotlib.pyplot as plt
```

```
[32]: #set current directory and create images directory
current_dir = Path(os.getcwd()).absolute()
images_dir = current_dir.joinpath('images')
images_dir.mkdir(parents=True, exist_ok = True)
```

```
[71]: #Used article "Simple Image Classification with ResNet-50" from Medium
#I know I can loop this, but just having trouble getting it to work
#Load image1 and set target to 224, 224 since that is the format ResNet50 needs
img1 = image.load_img('images/bunny.jpg', target_size = (224, 224))
#show image
plt.imshow(img1)
```

```
[71]: <matplotlib.image.AxesImage at 0x7fb08c25f640>
```



```
[72]: #turn into a numpy array
img1 = image.img_to_array(img1)
#insert new axis that will appear at the axis position in expanded array shape
img1 = np.expand_dims(img1, axis = 0)
#Preprocessing the numpy array encoding a batch of images
img1 = preprocess_input(img1)
#instantiating the ResNet50 model
model = ResNet50(weights = 'imagenet')
#predicting on the model and printing the result
preds1 = model.predict(img1)
print('Predicted:', decode_predictions(preds1, top = 1)[0])
```

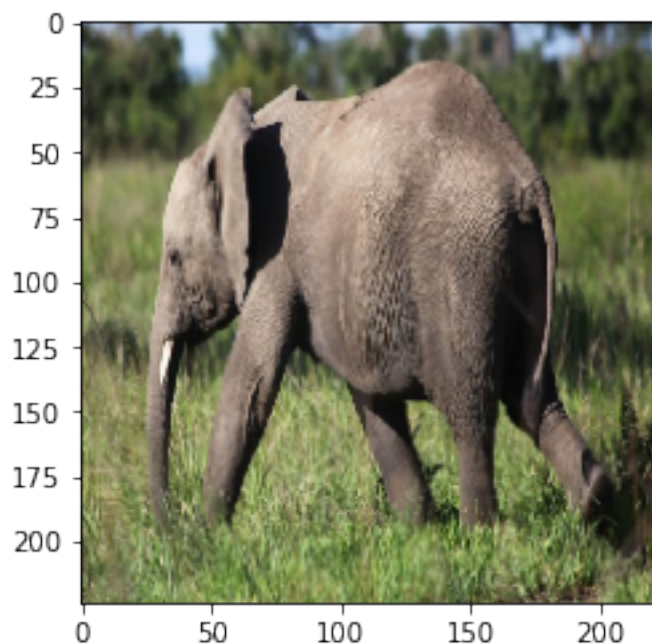
WARNING:tensorflow:5 out of the last 9 calls to <function Model.make_predict_function.<locals>.predict_function at 0x7fb0206641f0> triggered tf.function retracing. Tracing is expensive and the excessive number of tracings could be due to (1) creating @tf.function repeatedly in a loop, (2) passing tensors with different shapes, (3) passing Python objects instead of

tensors. For (1), please define your @tf.function outside of the loop. For (2), @tf.function has experimental_relax_shapes=True option that relaxes argument shapes that can avoid unnecessary retracing. For (3), please refer to https://www.tensorflow.org/guide/function#controlling_retracing and https://www.tensorflow.org/api_docs/python/tf/function for more details.

Predicted: [('n02326432', 'hare', 0.5823781)]

```
[73]: #Load image2 and set target to 224, 224 since that is the format ResNet50 needs  
img2 = image.load_img('images/elephant.jpg', target_size = (224, 224))  
#show image  
plt.imshow(img2)
```

[73]: <matplotlib.image.AxesImage at 0x7fb02052e5e0>

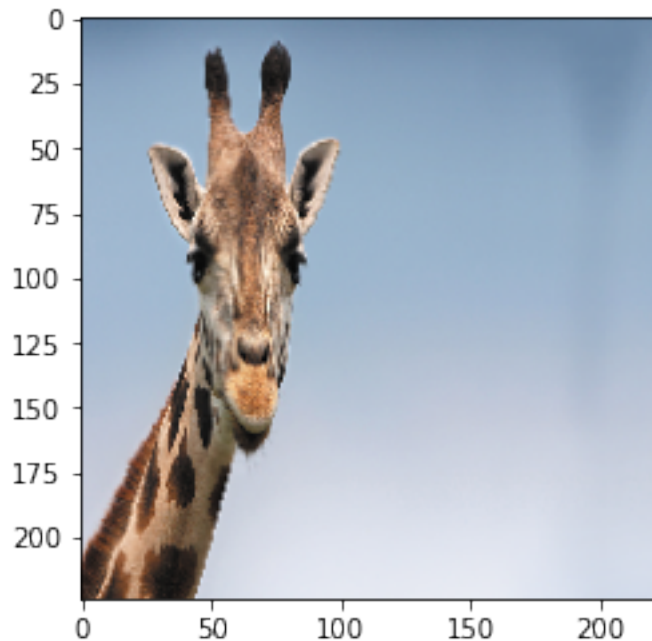


```
[74]: #turn into a numpy array  
img2 = image.img_to_array(img2)  
#insert new axis that will appear at the axis position in expanded array shape  
img2 = np.expand_dims(img2, axis = 0)  
#Preprocessing the numpy array encoding a batch of images  
img2 = preprocess_input(img2)  
#instantiating the ResNet50 model  
model2 = ResNet50(weights = 'imagenet')  
#predicting on the model and printing the result  
preds2 = model.predict(img2)  
print('Predicted:', decode_predictions(preds2, top = 1)[0])
```

Predicted: [('n01871265', 'tusker', 0.45374754)]

```
[75]: #Load image3 and set target to 224, 224 since that is the format ResNet50 needs
img3 = image.load_img('images/giraffe.png', target_size = (224, 224))
#show image
plt.imshow(img3)
```

[75]: <matplotlib.image.AxesImage at 0x7fb01855efd0>



```
[76]: #turn into a numpy array
img3 = image.img_to_array(img3)
#insert new axis that will appear at the axis position in expanded array shape
img3 = np.expand_dims(img3, axis = 0)
#Preprocessing the numpy array encoding a batch of images
img3 = preprocess_input(img3)
#instantiating the ResNet50 model
model3 = ResNet50(weights = 'imagenet')
#predicting on the model and printing the result
preds3 = model3.predict(img3)
print('Predicted:', decode_predictions(preds3, top = 1)[0])
```

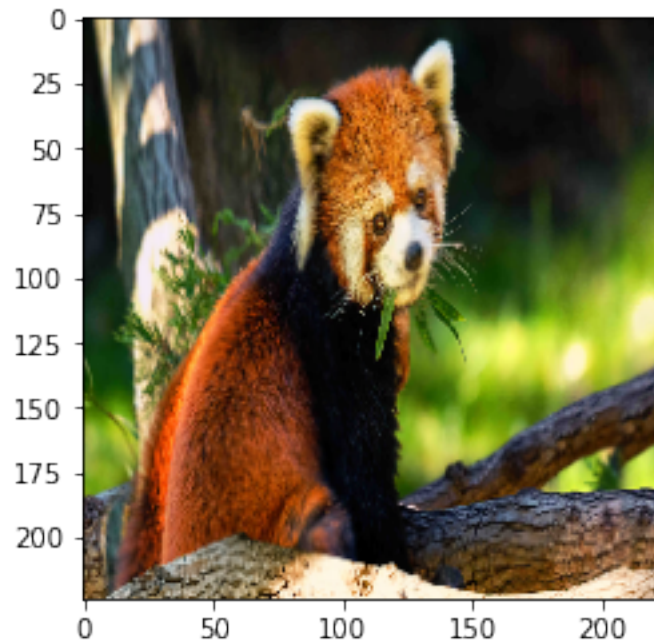
Predicted: [('n02417914', 'ibex', 0.36325175)]

```
[77]: #Load image4 and set target to 224, 224 since that is the format ResNet50 needs
img4 = image.load_img('images/red_panda.png', target_size = (224, 224))
#show image
```



```
plt.imshow(img4)
```

[77]: <matplotlib.image.AxesImage at 0x7fb0182d3940>




```
[78]: #turn into a numpy array
img4 = image.img_to_array(img4)
#insert new axis that will appear at the axis position in expanded array shape
img4 = np.expand_dims(img4, axis = 0)
#Preprocessing the numpy array encoding a batch of images
img4 = preprocess_input(img4)
#instantiating the ResNet50 model
model4 = ResNet50(weights = 'imagenet')
#predicting on the model and printing the result
preds4 = model.predict(img4)
print('Predicted:', decode_predictions(preds4, top = 1)[0])
```

Predicted: [('n02509815', 'lesser_panda', 0.977285)]

```
[79]: #Load image5 and set target to 224, 224 since that is the format ResNet50 needs
img5 = image.load_img('images/tiger.jpeg', target_size = (224, 224))
#show image
plt.imshow(img5)
```

[79]: <matplotlib.image.AxesImage at 0x7fb08c682e20>



```
[80]: #turn into a numpy array
img5 = image.img_to_array(img5)
#insert new axis that will appear at the axis position in expanded array shape
img5 = np.expand_dims(img5, axis = 0)
#Preprocessing the numpy array encoding a batch of images
img5 = preprocess_input(img5)
#instantiating the ResNet50 model
model5 = ResNet50(weights = 'imagenet')
#predicting on the model and printing the result
preds5 = model5.predict(img5)
print('Predicted:', decode_predictions(preds5, top = 1)[0])
```

Predicted: [('n02129604', 'tiger', 0.7966444)]