

assignment12_FoxAndrea

May 30, 2021

Name: Andrea Fox
Date: May 29, 2021
Course: DSC650 - T301 Big Data
Assignment 12

0.1 Assignment 12

Using section 8.4 in Deep Learning with Python as a guide, implement a variational autoencoder

```
[8]: #load libraries
import keras
from keras import layers
import tensorflow.compat.v1.keras.backend as K
from tensorflow.keras.models import Model
import numpy as np
from tensorflow.keras.datasets import mnist
import tensorflow as tf
import os
from pathlib import Path
import matplotlib.pyplot as plt
from scipy.stats import norm
```

```
[16]: #set directories and create results folder
current_dir = Path(os.getcwd()).absolute()
results_dir = current_dir.joinpath('results')
results_dir.mkdir(parents=True, exist_ok=True)
vae_dir = results_dir.joinpath('vae')

# Fixes version incompatibility errors between Keras and Tensorflow (saw this_
↪from Robert on general chat)
tf.compat.v1.disable_eager_execution()
```

```
[10]: #VAE encoder network
img_shape = (28, 28, 1)
batch_size = 16
latent_dim = 2 # Dimensionality of the latent space: a plane

input_img = keras.Input(shape=img_shape)
```

```

x = layers.Conv2D(32, 3,
                  padding='same', activation='relu')(input_img)
x = layers.Conv2D(64, 3,
                  padding='same', activation='relu',
                  strides=(2, 2))(x)
x = layers.Conv2D(64, 3,
                  padding='same', activation='relu')(x)
x = layers.Conv2D(64, 3,
                  padding='same', activation='relu')(x)
shape_before_flattening = K.int_shape(x)

x = layers.Flatten()(x)
x = layers.Dense(32, activation='relu')(x)

z_mean = layers.Dense(latent_dim)(x)
z_log_var = layers.Dense(latent_dim)(x)

```

```

[11]: #latent-space-sampling function
def sampling(args):
    z_mean, z_log_var = args
    epsilon = K.random_normal(shape=(K.shape(z_mean)[0], latent_dim),
                              mean=0., stddev=1.)
    return z_mean + K.exp(z_log_var) * epsilon

z = layers.Lambda(sampling)([z_mean, z_log_var])

```

```

[12]: #VAE encoder network, mapping latent space points to images
# This is the input where we will feed `z`.
decoder_input = layers.Input(K.int_shape(z)[1:])

# Upsample to the correct number of units
x = layers.Dense(np.prod(shape_before_flattening[1:]),
                  activation='relu')(decoder_input)

# Reshape into an image of the same shape as before our last `Flatten` layer
x = layers.Reshape(shape_before_flattening[1:])(x)

# We then apply then reverse operation to the initial
# stack of convolution layers: a `Conv2DTranspose` layers
# with corresponding parameters.
x = layers.Conv2DTranspose(32, 3,
                           padding='same', activation='relu',
                           strides=(2, 2))(x)
x = layers.Conv2D(1, 3,
                  padding='same', activation='sigmoid')(x)
# We end up with a feature map of the same size as the original input.

```

```

# This is our decoder model.
decoder = Model(decoder_input, x)

# We then apply it to `z` to recover the decoded `z`.
z_decoded = decoder(z)

```

```

[13]: #Custom layer used to compute the VAE loss
class CustomVariationalLayer(keras.layers.Layer):

    def vae_loss(self, x, z_decoded):
        x = K.flatten(x)
        z_decoded = K.flatten(z_decoded)
        xent_loss = keras.metrics.binary_crossentropy(x, z_decoded)
        kl_loss = -5e-4 * K.mean(
            1 + z_log_var - K.square(z_mean) - K.exp(z_log_var), axis=-1)
        return K.mean(xent_loss + kl_loss)

    def call(self, inputs):
        x = inputs[0]
        z_decoded = inputs[1]
        loss = self.vae_loss(x, z_decoded)
        self.add_loss(loss, inputs=inputs)
        # We don't use this output.
        return x

# We call our custom layer on the input and the decoded output,
# to obtain the final model output.
y = CustomVariationalLayer()([input_img, z_decoded])

```

```

[14]: #training the VAE
vae = Model(input_img, y)
vae.compile(optimizer='rmsprop', loss=None)
vae.summary()

# Train the VAE on MNIST digits
(x_train, _), (x_test, y_test) = mnist.load_data()

x_train = x_train.astype('float32') / 255.
x_train = x_train.reshape(x_train.shape + (1,))
x_test = x_test.astype('float32') / 255.
x_test = x_test.reshape(x_test.shape + (1,))

vae.fit(x=x_train, y=None,
        shuffle=True,
        epochs=10,
        batch_size=batch_size,
        validation_data=(x_test, None))

```

WARNING:tensorflow:Output custom_variational_layer_1 missing from loss dictionary. We assume this was done on purpose. The fit and evaluate APIs will not be expecting any data to be passed to custom_variational_layer_1.

Model: "model_3"

Layer (type)	Output Shape	Param #	Connected to
=====			
input_1 (InputLayer)	[(None, 28, 28, 1)]	0	

conv2d_5 (Conv2D)	(None, 28, 28, 32)	320	input_1[0][0]

conv2d_6 (Conv2D)	(None, 14, 14, 64)	18496	conv2d_5[0][0]

conv2d_7 (Conv2D)	(None, 14, 14, 64)	36928	conv2d_6[0][0]

conv2d_8 (Conv2D)	(None, 14, 14, 64)	36928	conv2d_7[0][0]

flatten_1 (Flatten)	(None, 12544)	0	conv2d_8[0][0]

dense_4 (Dense)	(None, 32)	401440	flatten_1[0][0]

dense_5 (Dense)	(None, 2)	66	dense_4[0][0]

dense_6 (Dense)	(None, 2)	66	dense_4[0][0]

lambda_1 (Lambda)	(None, 2)	0	dense_5[0][0] dense_6[0][0]

model_2 (Functional)	(None, 28, 28, 1)	56385	lambda_1[0][0]

custom_variational_layer_1 (Cus	(None, 28, 28, 1)	0	input_1[0][0] model_2[0][0]
=====			

Total params: 550,629

Trainable params: 550,629
Non-trainable params: 0

```
-----  
-----  
Train on 60000 samples, validate on 10000 samples  
Epoch 1/10  
59968/60000 [=====>.] - ETA: 0s - loss: 0.2334  
  
/opt/conda/lib/python3.8/site-  
packages/tensorflow/python/keras/engine/training.py:2325: UserWarning:  
`Model.state_updates` will be removed in a future version. This property should  
not be used in TensorFlow 2.0, as `updates` are applied automatically.  
    warnings.warn("`Model.state_updates` will be removed in a future version. '  
  
60000/60000 [=====] - 85s 1ms/sample - loss: 0.2333 -  
val_loss: 0.1987  
Epoch 2/10  
60000/60000 [=====] - 84s 1ms/sample - loss: 0.1952 -  
val_loss: 0.1920  
Epoch 3/10  
60000/60000 [=====] - 83s 1ms/sample - loss: 0.1906 -  
val_loss: 0.1880  
Epoch 4/10  
60000/60000 [=====] - 83s 1ms/sample - loss: 0.1878 -  
val_loss: 0.1869  
Epoch 5/10  
60000/60000 [=====] - 83s 1ms/sample - loss: 0.1859 -  
val_loss: 0.1846  
Epoch 6/10  
60000/60000 [=====] - 82s 1ms/sample - loss: 0.1845 -  
val_loss: 0.1835  
Epoch 7/10  
60000/60000 [=====] - 83s 1ms/sample - loss: 0.1832 -  
val_loss: 0.1880  
Epoch 8/10  
60000/60000 [=====] - 83s 1ms/sample - loss: 0.1823 -  
val_loss: 0.1815  
Epoch 9/10  
60000/60000 [=====] - 83s 1ms/sample - loss: 0.1816 -  
val_loss: 0.1825  
Epoch 10/10  
60000/60000 [=====] - 83s 1ms/sample - loss: 0.1809 -  
val_loss: 0.1824
```

[14]: <tensorflow.python.keras.callbacks.History at 0x7f7738110fd0>

[17]: *#Sampling a grid of points from the 2D latent space and decoding them to images
Display a 2D manifold of the digits*

```

n = 15 # figure with 15x15 digits
digit_size = 28
figure = np.zeros((digit_size * n, digit_size * n))
# Linearly spaced coordinates on the unit square were transformed
# through the inverse CDF (ppf) of the Gaussian
# to produce values of the latent variables z,
# since the prior of the latent space is Gaussian
grid_x = norm.ppf(np.linspace(0.05, 0.95, n))
grid_y = norm.ppf(np.linspace(0.05, 0.95, n))

for i, yi in enumerate(grid_x):
    for j, xi in enumerate(grid_y):
        z_sample = np.array([[xi, yi]])
        z_sample = np.tile(z_sample, batch_size).reshape(batch_size, 2)
        x_decoded = decoder.predict(z_sample, batch_size=batch_size)
        digit = x_decoded[0].reshape(digit_size, digit_size)
        figure[i * digit_size: (i + 1) * digit_size,
              j * digit_size: (j + 1) * digit_size] = digit

plt.figure(figsize=(10, 10))
plt.imshow(figure, cmap='Greys_r')

#Robert mentioned this in general chat that he used so it saved it as well as
↳ showed in output
plt.savefig(vae_dir.joinpath('15x15_Grid_of_Var_Encod_MNIST_Nums.png'),
↳ bbox_inches='tight')

```

